

# sta210 project

Cole Walker, Madison Griffin

## loading packages & dataset

```
library(tidyverse)
library(tidymodels)
library(readxl)
library(MASS)
library(leaps)
library(caret)
library(glmnet)
library(Stat2Data)
#library(statnnet)
library(lme4)
library(UpSetR)
library(nlme)
library(sjstats)
set.seed(8)
soccer <- read_excel("AllTimeRankingByClub.xlsx")
```

## Introduction and data

## Data Cleaning

```
soccer = soccer %>%
  rename(goals_for = `Goals For`,
         goals_against = `Goals Against`,
         goal_diff = `Goal Diff`)

soccer = soccer %>%
  mutate(winspermatch = Win/Played,
```

```

    pointspermatch = Pts/Played,
    goalspermatch = goals_for/Played,
    goalsagainstpermatch = goals_against/Played,
    goalmarginpermatch = goal_diff/Played)

soccer = soccer %>%
  mutate(topfiveleague = ifelse(Country == 'ESP' | Country == 'ENG' | Country == 'GER' | C

```

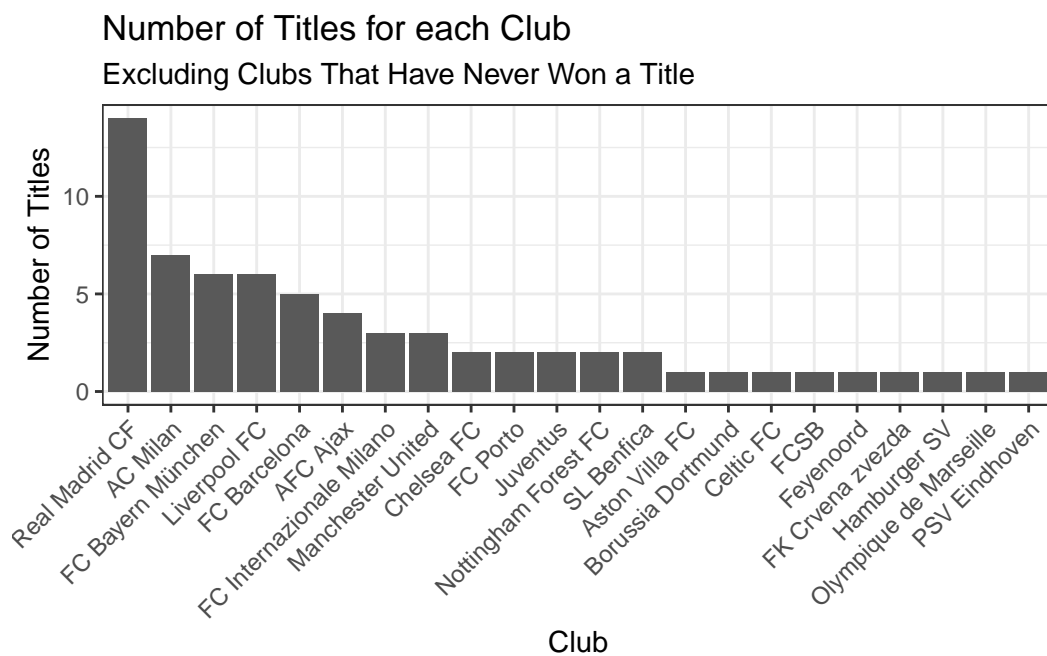
## EDA

Plot 1:

```

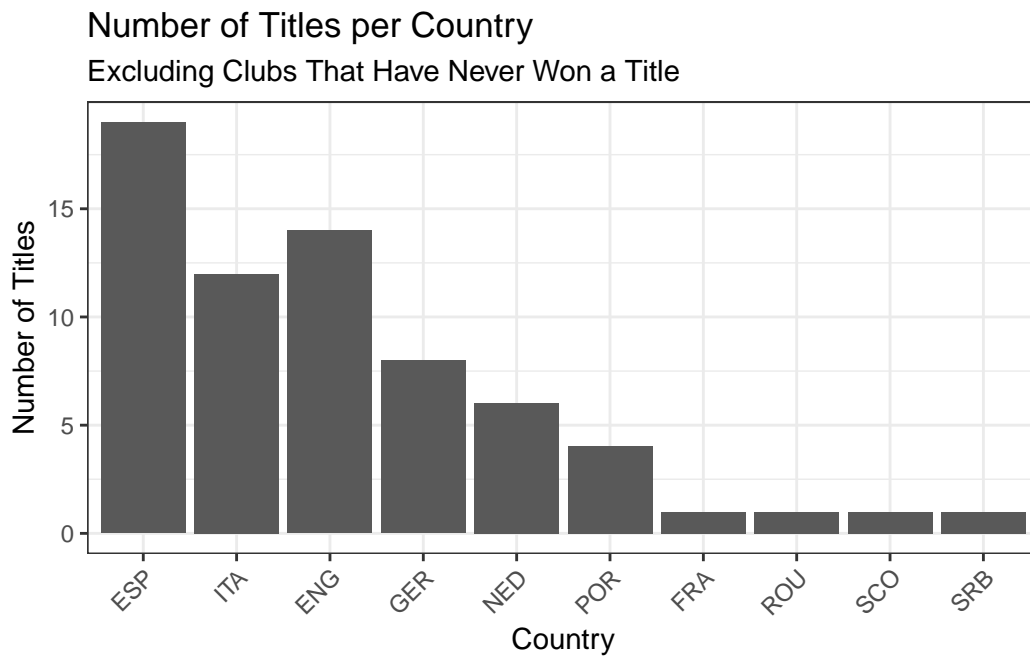
soccer %>%
  filter(Titles > 0) %>%
  ggplot(aes(x = reorder(Club, (-Titles)), y = Titles)) +
  geom_bar(stat = 'identity') +
  labs(x = 'Club', y = 'Number of Titles', title = 'Number of Titles for each Club',
       subtitle = 'Excluding Clubs That Have Never Won a Title') +
  theme_bw() +
  scale_x_discrete(guide = guide_axis(angle = 45))

```



Plot 2:

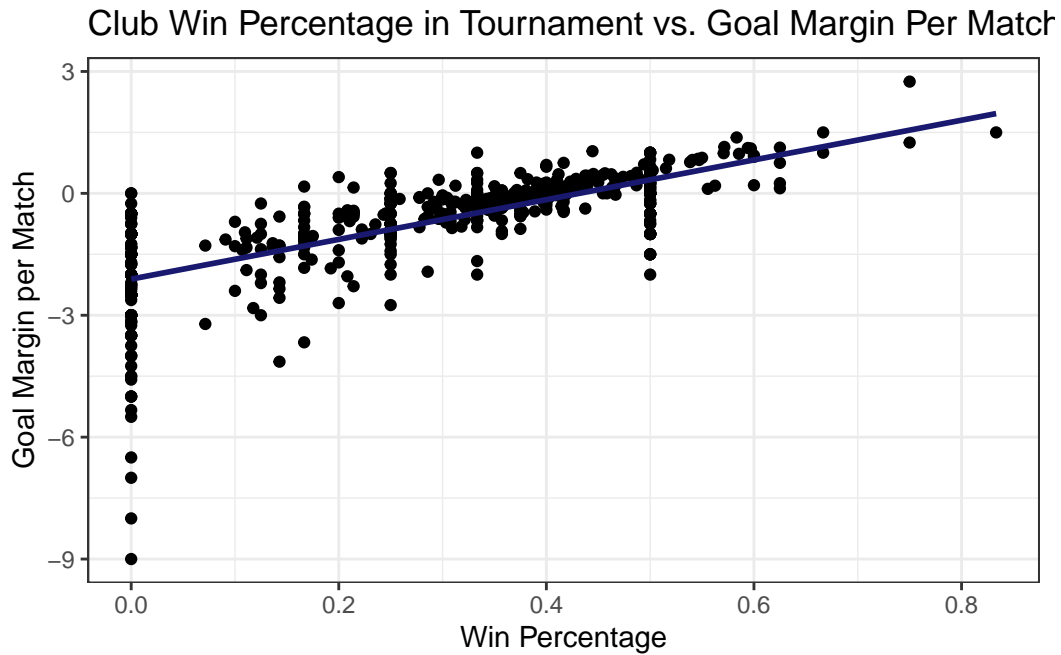
```
soccer %>%  
  filter(Titles > 0) %>%  
  ggplot(aes(x = reorder(Country, (-Titles)), y = Titles)) +  
  geom_bar(stat = 'identity') +  
  labs(x = 'Country', y = 'Number of Titles', title = 'Number of Titles per Country',  
       subtitle = 'Excluding Clubs That Have Never Won a Title') +  
  theme_bw() +  
  scale_x_discrete(guide = guide_axis(angle = 45))
```



Plot 3:

```
ggplot(data = soccer, aes(x = winspermatch, y = goalmarginpermatch)) +  
  geom_point() +  
  geom_smooth(method = 'lm', se = F, color = 'midnightblue') +  
  labs(x = "Win Percentage", y = 'Goal Margin per Match',  
       title = 'Club Win Percentage in Tournament vs. Goal Margin Per Match') +  
  theme_bw()
```

`geom\_smooth()` using formula = 'y ~ x'



## Methods

Model 1: Linear Regression

Outcome:

- points per match

Predictors:

- wins per match
- goals per match
- goals against per match
- goal margin per match
- top five league

```
model1 = lm(pointspermatch ~ winspermatch + goalspermatch +  
             goalsagainstpermatch + goalmarginpermatch + topfiveleague,  
             data = soccer)  
summary(model1)
```

Call:

```
lm(formula = pointspermatch ~ winspermatch + goalspermatch +  
    goalsagainstpermatch + goalmarginpermatch + topfiveleague,  
    data = soccer)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.33721	-0.06803	0.00581	0.06882	0.56561

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.435103	0.020395	21.334	< 2e-16 ***
winspermatch	1.381209	0.049182	28.084	< 2e-16 ***
goalspermatch	0.096467	0.014703	6.561	1.28e-10 ***
goalsagainstpermatch	-0.097891	0.006649	-14.723	< 2e-16 ***
goalmarginpermatch	NA	NA	NA	NA
topfiveleaguertop five	0.025705	0.016498	1.558	0.12

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1355 on 525 degrees of freedom

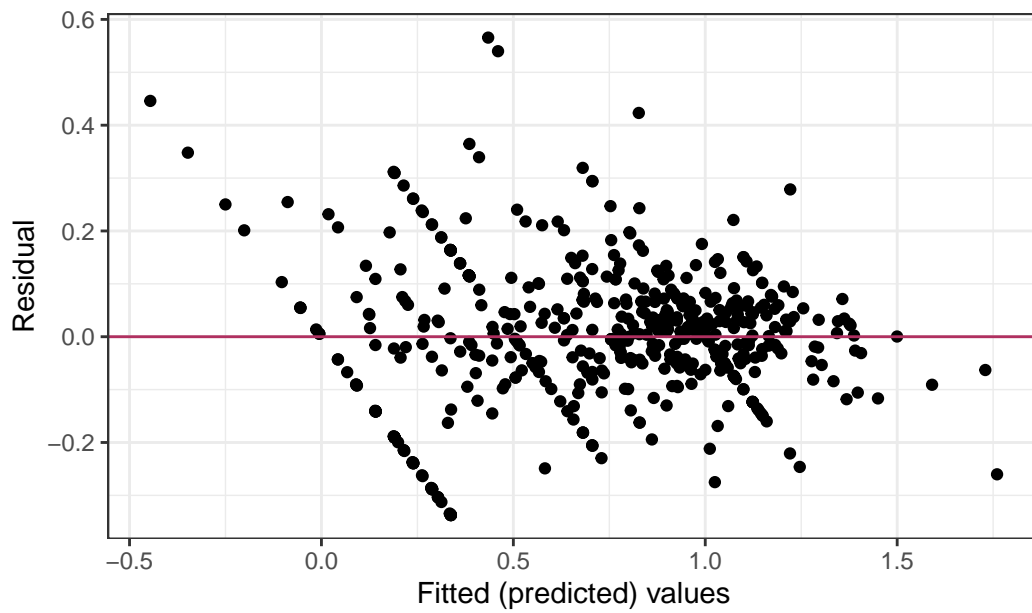
Multiple R-squared: 0.8845, Adjusted R-squared: 0.8836

F-statistic: 1005 on 4 and 525 DF, p-value: < 2.2e-16

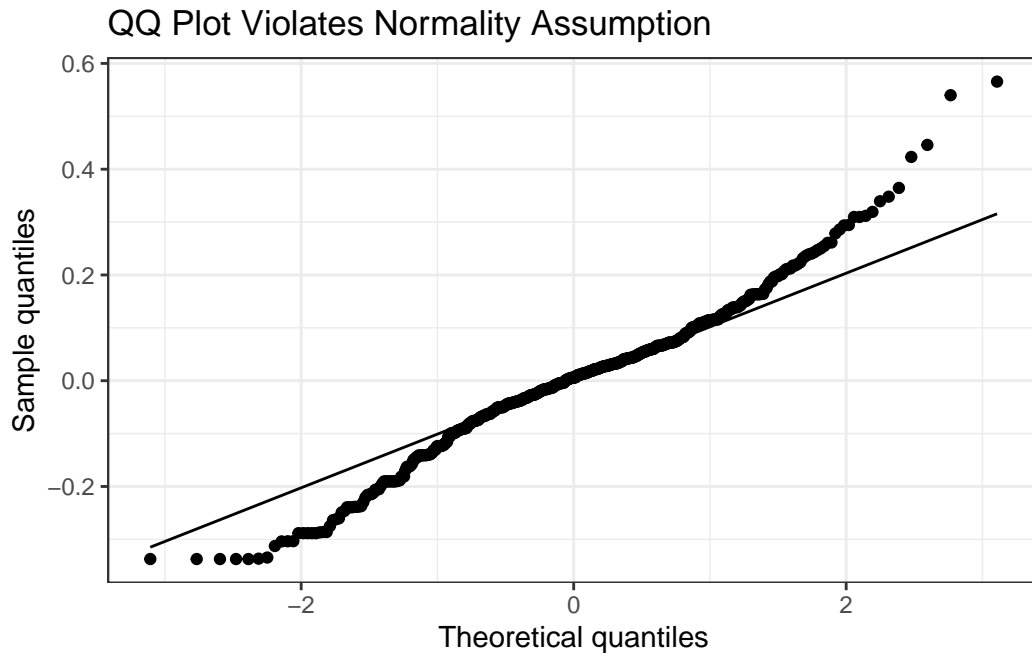
Conditions for Model 1: Violated linearity, constant variance, and normality

```
modell1aug = augment(modell1)  
  
ggplot(modell1aug, aes(x = .fitted, y = .resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, color = 'maroon') +  
  labs(x = "Fitted (predicted) values", y = 'Residual') +  
  ggtitle('Residual Plot Violates Linearity & Constant Variance Assumptions') +  
  theme_bw()
```

Residual Plot Violates Linearity & Constant Variance Assumpt



```
ggplot(modell1aug, aes(sample = .resid)) +  
  stat_qq() +  
  stat_qq_line() +  
  theme_bw() +  
  labs(x = 'Theoretical quantiles',  
       y = 'Sample quantiles',  
       title = 'QQ Plot Violates Normality Assumption')
```



testing correlations because something is weird

```
test = soccer %>%
  dplyr::select(pointspermatch, winspermatch, goalspermatch, goalsagainstpermatch, goalmarginpermatch)

cor(test)
```

	pointspermatch	winspermatch	goalspermatch
pointspermatch	1.0000000	0.9103634	0.7085198
winspermatch	0.9103634	1.0000000	0.7213852
goalspermatch	0.7085198	0.7213852	1.0000000
goalsagainstpermatch	-0.6382465	-0.4996595	-0.2858228
goalmarginpermatch	0.8108285	0.7085429	0.6652299

	goalsagainstpermatch	goalmarginpermatch
pointspermatch	-0.6382465	0.8108285
winspermatch	-0.4996595	0.7085429
goalspermatch	-0.2858228	0.6652299
goalsagainstpermatch	1.0000000	-0.9056286
goalmarginpermatch	-0.9056286	1.0000000

Model 2: Linear Regression

Outcome:

- points per match

Predictors:

- wins per match
- goals per match
- goals against per match
- top five league

```
model2 = lm(pointspermatch ~ winspermatch + goalspermatch +
             goalsagainstpermatch + topfiveleague,
             data = soccer)
summary(model2)
```

Call:

```
lm(formula = pointspermatch ~ winspermatch + goalspermatch +
    goalsagainstpermatch + topfiveleague, data = soccer)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.33721	-0.06803	0.00581	0.06882	0.56561

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.435103	0.020395	21.334	< 2e-16 ***
winspermatch	1.381209	0.049182	28.084	< 2e-16 ***
goalspermatch	0.096467	0.014703	6.561	1.28e-10 ***
goalsagainstpermatch	-0.097891	0.006649	-14.723	< 2e-16 ***
topfiveleague	0.025705	0.016498	1.558	0.12

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1355 on 525 degrees of freedom

Multiple R-squared: 0.8845, Adjusted R-squared: 0.8836

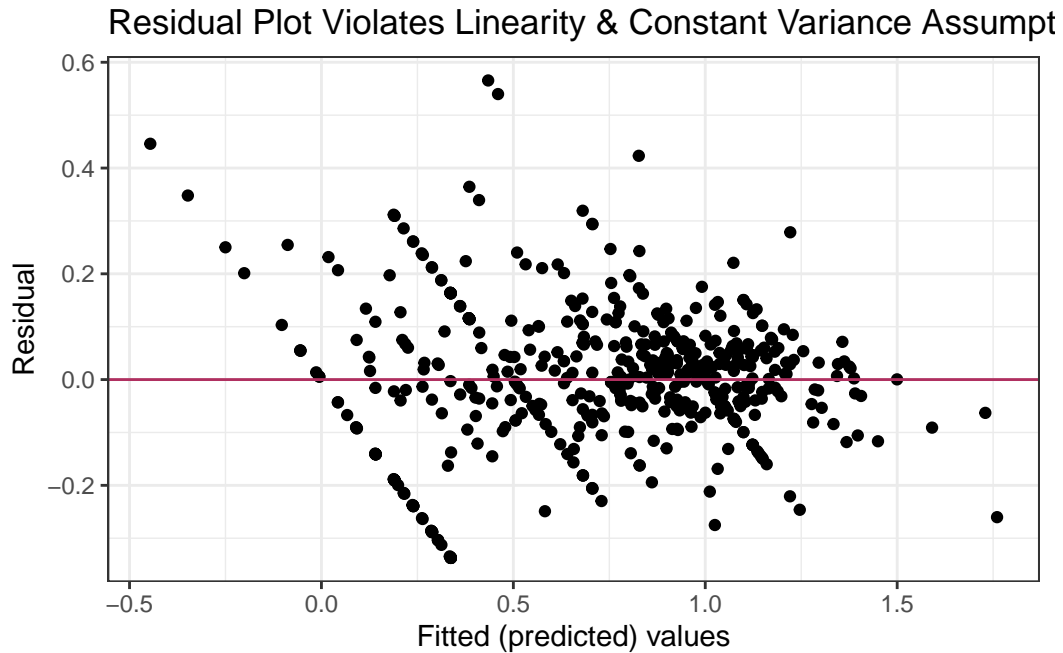
F-statistic: 1005 on 4 and 525 DF, p-value: < 2.2e-16

Check Assumptions for Model 2 still all violated

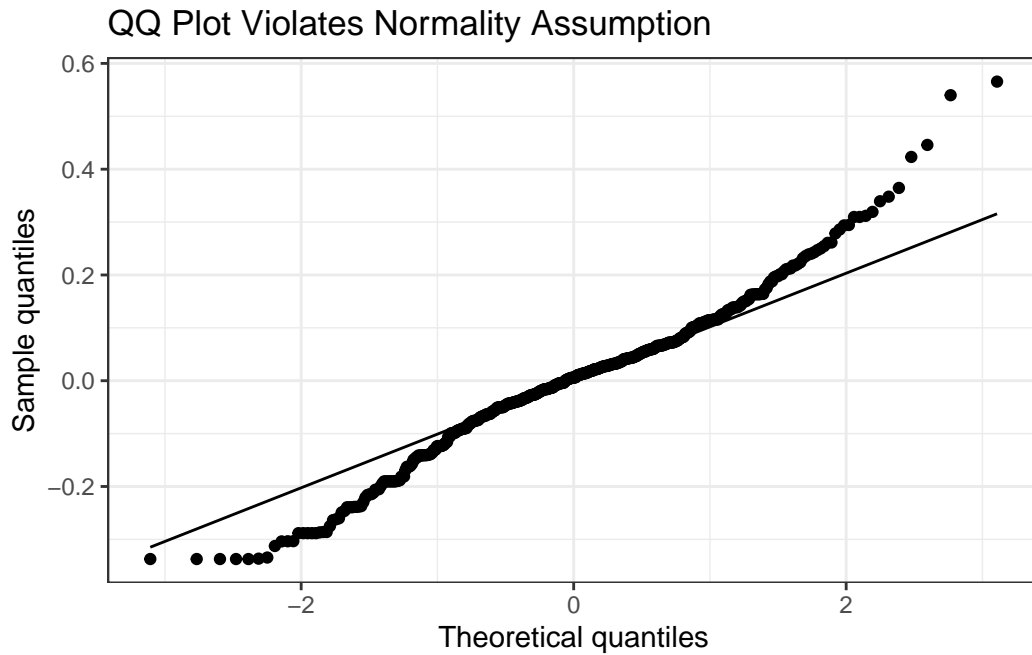


```
model2aug = augment(model2)
```

```
ggplot(model2aug, aes(x = .fitted, y = .resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, color = 'maroon') +  
  labs(x = "Fitted (predicted) values", y = 'Residual') +  
  ggtitle('Residual Plot Violates Linearity & Constant Variance Assumptions') +  
  theme_bw()
```



```
ggplot(model2aug, aes(sample = .resid)) +  
  stat_qq() +  
  stat_qq_line() +  
  theme_bw() +  
  labs(x = 'Theoretical quantiles',  
       y = 'Sample quantiles',  
       title = 'QQ Plot Violates Normality Assumption')
```



Model 3: Linear Regression

Outcome:

- points per match

Predictors:

- wins per match
- goal margin per match
- top five league

```
model3 = lm(pointspermatch ~ winspermatch + goalmarginpermatch + topfiveleague,  
             data = soccer)  
summary(model3)
```

Call:

```
lm(formula = pointspermatch ~ winspermatch + goalmarginpermatch +  
    topfiveleague, data = soccer)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.33632	-0.06813	0.00603	0.06892	0.56598

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.434020	0.016755	25.904	<2e-16 ***
winspermatch	1.379068	0.043462	31.731	<2e-16 ***
goalmarginpermatch	0.097704	0.006335	15.423	<2e-16 ***
topfiveleaguertop five	0.025563	0.016412	1.558	0.12

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1354 on 526 degrees of freedom

Multiple R-squared: 0.8845, Adjusted R-squared: 0.8838

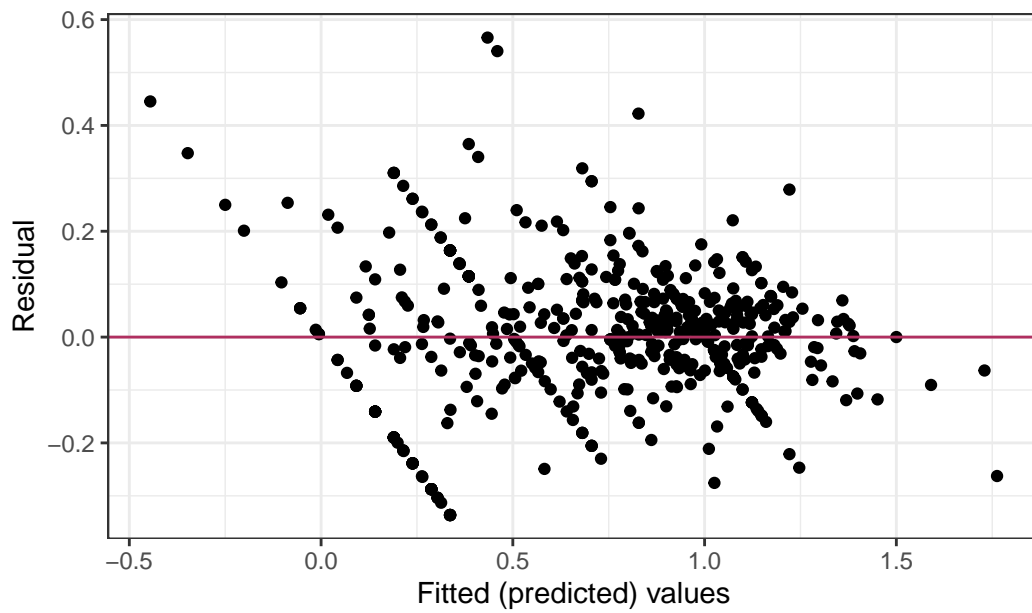
F-statistic: 1343 on 3 and 526 DF, p-value: < 2.2e-16

Check Assumptions Model 3

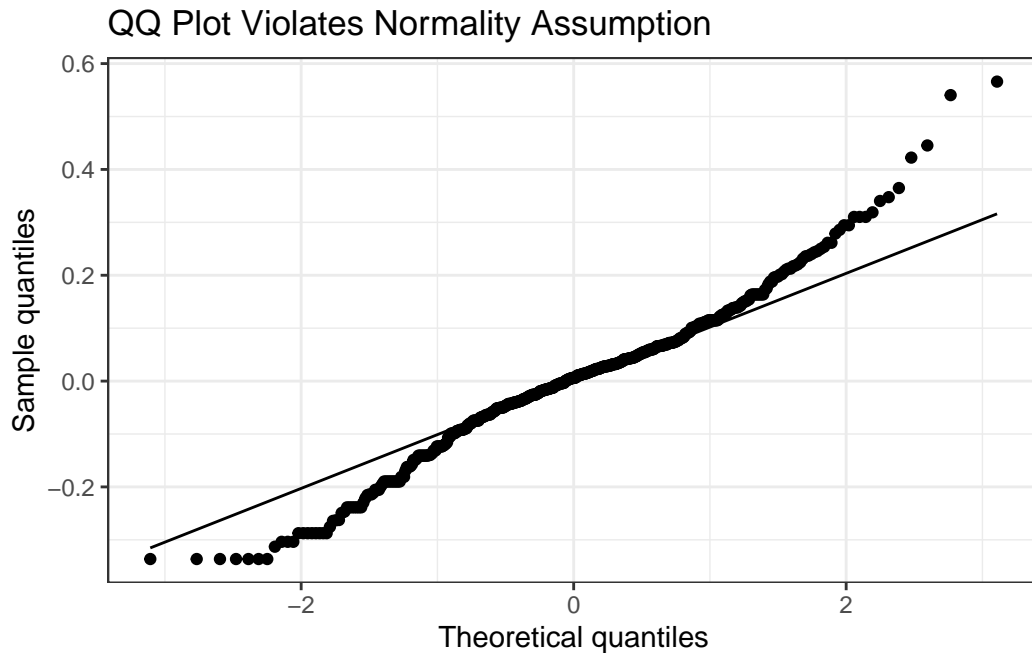
```
model3aug = augment(model3)

ggplot(model3aug, aes(x = .fitted, y = .resid)) +
  geom_point() +
  geom_hline(yintercept = 0, color = 'maroon') +
  labs(x = "Fitted (predicted) values", y = 'Residual') +
  ggtitle('Residual Plot Violates Linearity & Constant Variance Assumptions') +
  theme_bw()
```

Residual Plot Violates Linearity & Constant Variance Assumpt



```
ggplot(model3aug, aes(sample = .resid)) +  
  stat_qq() +  
  stat_qq_line() +  
  theme_bw() +  
  labs(x = 'Theoretical quantiles',  
       y = 'Sample quantiles',  
       title = 'QQ Plot Violates Normality Assumption')
```



Model 4: Linear Mixed Effects Model (potential violation of independence with topfive-league)

Outcome:

- points per match

Predictors:

- random intercept for topfiveleague
- goal margin per match
- wins per match
- goals per match

```
model4 = lmer(pointspersmatch ~ 1 + goalspersmatch + winspersmatch + goalsagainstpersmatch +
              (1|topfiveleague), data = soccer)
summary(model4)
```

Linear mixed model fit by REML ['lmerMod']

Formula:

pointspersmatch ~ 1 + goalspersmatch + winspersmatch + goalsagainstpersmatch +

```
(1 | topfiveleague)
Data: soccer
```

REML criterion at convergence: -589.3

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.4908	-0.5166	0.0546	0.5113	4.1647

Random effects:

Groups	Name	Variance	Std.Dev.
topfiveleague	(Intercept)	0.0001943	0.01394
	Residual	0.0183620	0.13551

Number of obs: 530, groups: topfiveleague, 2

Fixed effects:

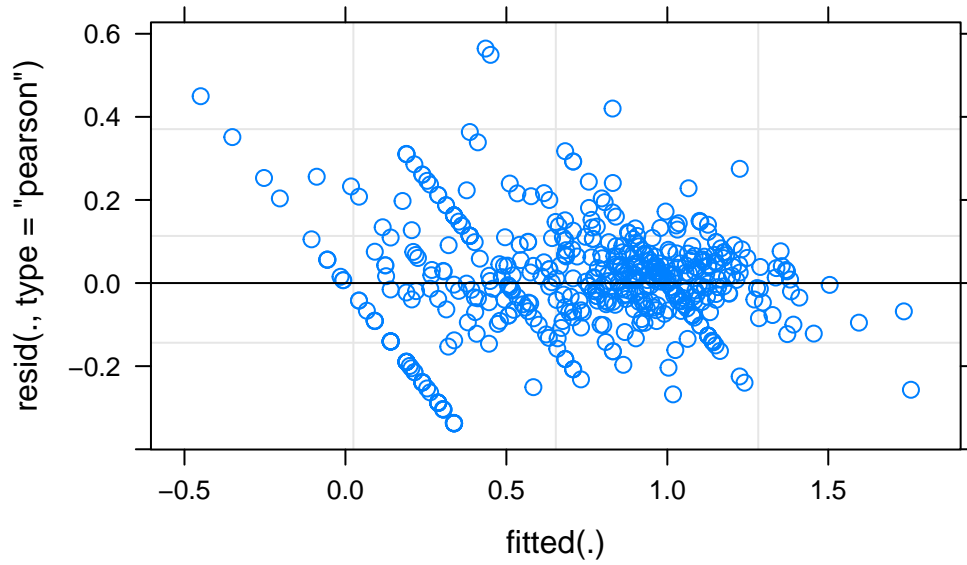
	Estimate	Std. Error	t value
(Intercept)	0.443476	0.023240	19.082
goalspermatch	0.097878	0.014635	6.688
winspermatch	1.383025	0.049148	28.140
goalsagainstpermatch	-0.098400	0.006629	-14.843

Correlation of Fixed Effects:

	(Intr)	glsprm	wnsprm
goalsprmtch	-0.237		
winspermtch	-0.342	-0.687	
glsgnstprmt	-0.676	-0.134	0.436

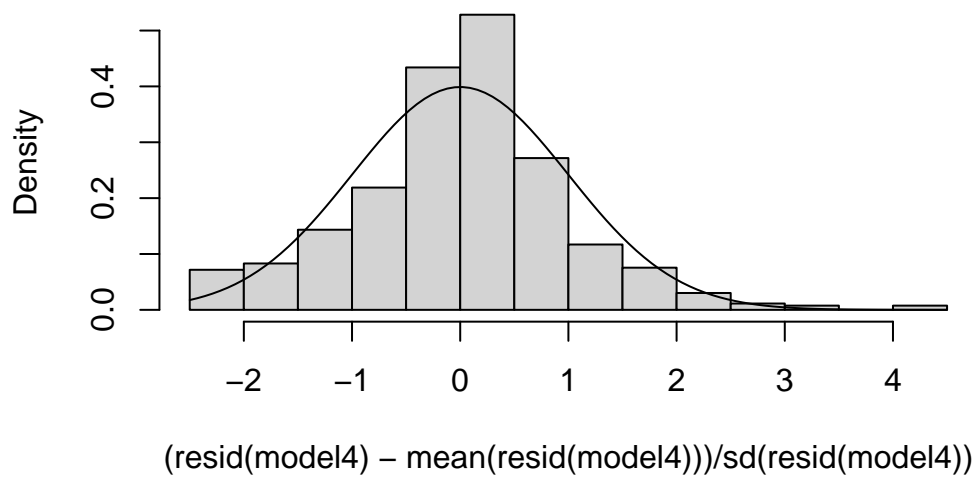
Checking assumptions model 4

```
plot(model4)
```



```
hist((resid(model4) - mean(resid(model4))) / sd(resid(model4))), freq = FALSE); curve(dnorm
```

istogram of  $(\text{resid}(\text{model4}) - \text{mean}(\text{resid}(\text{model4}))) / \text{sd}(\text{resid}(\text{model4}))$



## Model 5: Linear Regression

Outcome:

- points per match

Predictors:

- topfiveleague
- goal margin per match
- goals per match
- goals against per match

```
model5 = lm(pointspermatch ~ goalspermatch + goalsagainstpermatch + goalmarginpermatch +  
            topfiveleague, data = soccer)  
summary(model5)
```

Call:

```
lm(formula = pointspermatch ~ goalspermatch + goalsagainstpermatch +  
    goalmarginpermatch + topfiveleague, data = soccer)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.86804	-0.10210	0.01674	0.12150	0.95670

Coefficients: (1 not defined because of singularities)

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.649387	0.029890	21.726	<2e-16 ***
goalspermatch	0.377006	0.017049	22.113	<2e-16 ***
goalsagainstpermatch	-0.178454	0.009479	-18.826	<2e-16 ***
goalmarginpermatch	NA	NA	NA	NA
topfiveleague	0.052367	0.026029	2.012	0.0447 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2141 on 526 degrees of freedom

Multiple R-squared: 0.711, Adjusted R-squared: 0.7093

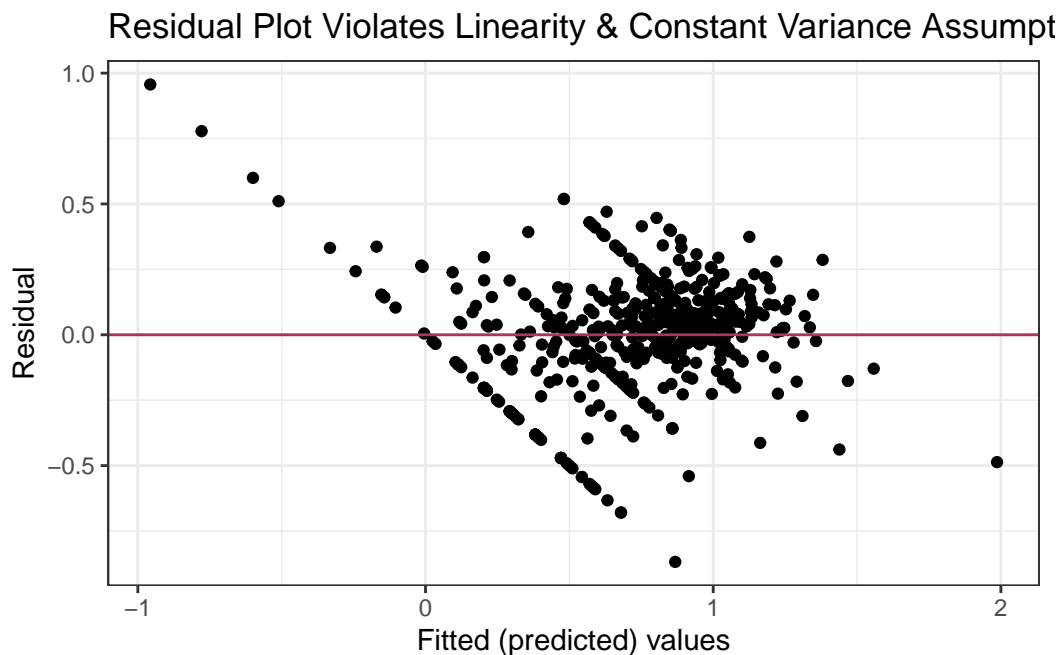
F-statistic: 431.3 on 3 and 526 DF, p-value: < 2.2e-16

Checking Assumptions

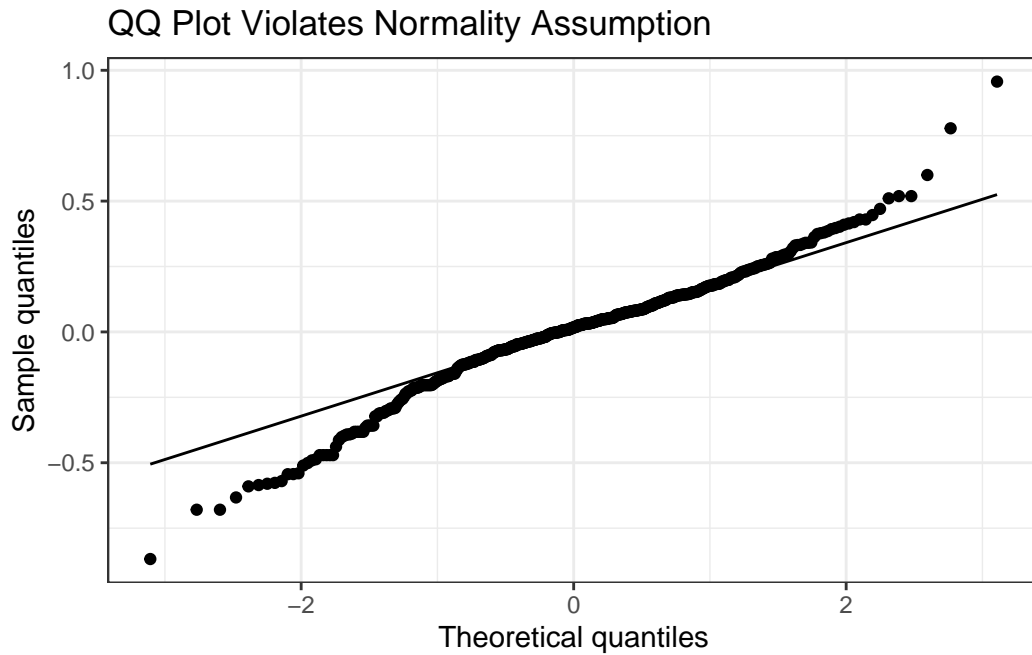


```
model5aug = augment(model5)
```

```
ggplot(model5aug, aes(x = .fitted, y = .resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, color = 'maroon') +  
  labs(x = "Fitted (predicted) values", y = 'Residual') +  
  ggtitle('Residual Plot Violates Linearity & Constant Variance Assumptions') +  
  theme_bw()
```



```
ggplot(model5aug, aes(sample = .resid)) +  
  stat_qq() +  
  stat_qq_line() +  
  theme_bw() +  
  labs(x = 'Theoretical quantiles',  
       y = 'Sample quantiles',  
       title = 'QQ Plot Violates Normality Assumption')
```



Model 6: linear mixed effects (random intercept for topfiveleague)

Outcome:

- points per match

Predictors:

- random intercept for topfiveleague
- goal margin per match
- goals per match

```
model6 = lmer(pointspermatch ~ 1 + goalspermatch + goalsagainstpermatch + goalmarginpermatch  
              (1|topfiveleague), data = soccer)
```

fixed-effect model matrix is rank deficient so dropping 1 column / coefficient

```
summary(model6)
```

```

Linear mixed model fit by REML ['lmerMod']
Formula:
pointspermatch ~ 1 + goalspermatch + goalsagainstpermatch + goalmarginpermatch +
  (1 | topfiveleague)
Data: soccer

```

REML criterion at convergence: -109.6

Scaled residuals:

Min	1Q	Median	3Q	Max
-4.0695	-0.4806	0.0845	0.5708	4.4929

Random effects:

Groups	Name	Variance	Std.Dev.
topfiveleague	(Intercept)	0.001032	0.03213
Residual		0.045860	0.21415

Number of obs: 530, groups: topfiveleague, 2

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.670447	0.038501	17.41
goalspermatch	0.379188	0.016909	22.43
goalsagainstpermatch	-0.179209	0.009449	-18.97

Correlation of Fixed Effects:

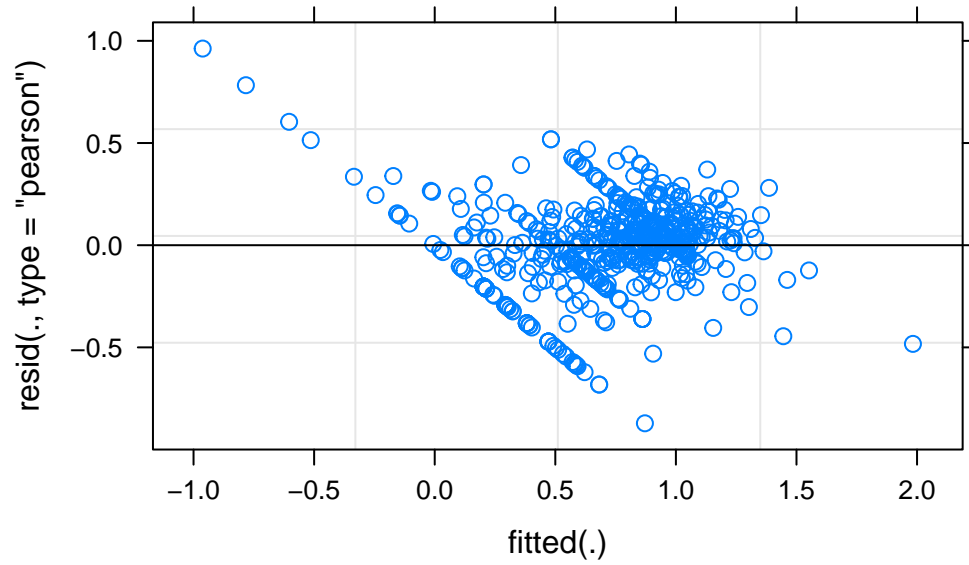
	(Intr)	glsprm
goalsprmtch	-0.627	
glsgnstprmt	-0.550	0.244

fit warnings:

fixed-effect model matrix is rank deficient so dropping 1 column / coefficient

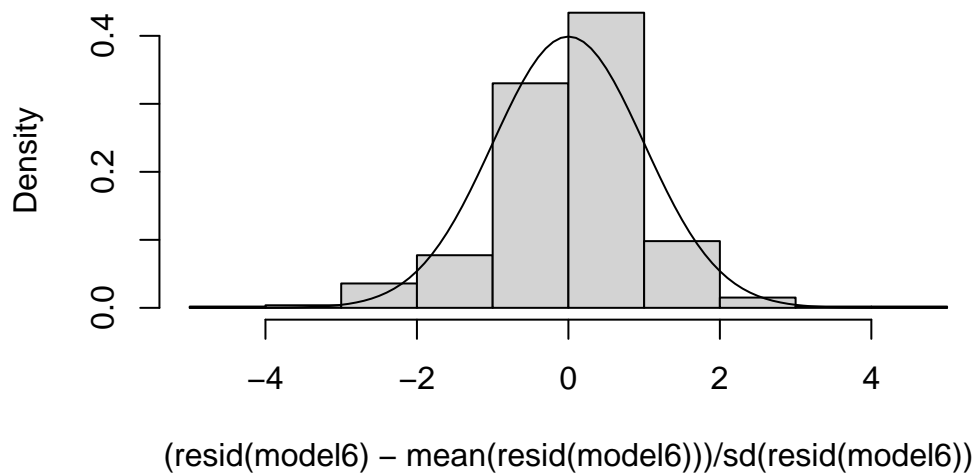
Checking assumptions:

```
plot(model6)
```



```
hist((resid(model6) - mean(resid(model6))) / sd(resid(model6)), freq = FALSE); curve(dnorm
```

istogram of  $(\text{resid}(\text{model6}) - \text{mean}(\text{resid}(\text{model6}))) / \text{sd}(\text{resid}(\text{model6}))$



## Variable Selection - LASSO

```
y = soccer$pointspermatch
x = model.matrix(pointspermatch ~ winspermatch + goalspermatch + goalsagainstpermatch +
                  goalmarginpermatch + topfiveleague, data = soccer)
m_lasso_cv = cv.glmnet(x, y, alpha = 1)

best_lambda = m_lasso_cv$lambda.min
best_lambda
```

```
[1] 0.003448411
```

```
m_best = glmnet(x, y, alpha = 1, lambda = best_lambda)
m_best$beta
```

```
6 x 1 sparse Matrix of class "dgCMatrix"
      s0
```

```
(Intercept)      .
winspermatch      1.37038572
goalspermatch     .
goalsagainstpermatch .
goalmarginpermatch 0.09659559
topfiveleague     0.01905888
```

```
bestlasso = lm(pointspermatch ~ winspermatch + goalmarginpermatch + topfiveleague,
                data = soccer)
```

## Variable Selection - Stepwise Selection

```
m_none = lm(pointspermatch ~ 1, data = soccer)
m_all = lm(pointspermatch ~ winspermatch + goalspermatch + goalsagainstpermatch +
            goalmarginpermatch + topfiveleague, data = soccer)
```

## Forward Selection

```
stepAIC(m_none,
        scope = list(lower = m_none, upper = m_all),
        data = soccer, direction = 'forward')
```

Start: AIC=-977.7  
 pointspermatch ~ 1

	Df	Sum of Sq	RSS	AIC
+ winspermatch	1	69.170	14.292	-1911.0
+ goalmarginpermatch	1	54.871	28.590	-1543.5
+ goalspermatch	1	41.898	41.564	-1345.2
+ goalsagainstpermatch	1	33.999	49.463	-1253.0
+ topfiveleague	1	9.140	74.322	-1037.2
<none>			83.462	-977.7

Step: AIC=-1910.99  
 pointspermatch ~ winspermatch

	Df	Sum of Sq	RSS	AIC
+ goalmarginpermatch	1	4.6072	9.6847	-2115.2
+ goalsagainstpermatch	1	3.7403	10.5516	-2069.8
+ goalspermatch	1	0.4669	13.8250	-1926.6
+ topfiveleague	1	0.2923	13.9996	-1919.9
<none>			14.2919	-1911.0

Step: AIC=-2115.24  
 pointspermatch ~ winspermatch + goalmarginpermatch

	Df	Sum of Sq	RSS	AIC
+ topfiveleague	1	0.044462	9.6402	-2115.7
<none>			9.6847	-2115.2
+ goalspermatch	1	0.000047	9.6846	-2113.2
+ goalsagainstpermatch	1	0.000047	9.6846	-2113.2

Step: AIC=-2115.68  
 pointspermatch ~ winspermatch + goalmarginpermatch + topfiveleague

	Df	Sum of Sq	RSS	AIC
<none>			9.6402	-2115.7
+ goalspermatch	1	0.00015982	9.6400	-2113.7
+ goalsagainstpermatch	1	0.00015982	9.6400	-2113.7

Call:  
 lm(formula = pointspermatch ~ winspermatch + goalmarginpermatch +  
 topfiveleague, data = soccer)

Coefficients:

(Intercept)	winspermatch	goalmarginpermatch
0.43402	1.37907	0.09770
topfiveleagu	top five	
0.02556		

```
bestforward = lm(pointspermatch ~ winspermatch + goalmarginpermatch +  
topfiveleague, data = soccer)
```

## Backward Selection

```
stepAIC(m_all,  
scope = list(lower = m_none, upper = m_all),  
data = soccer, direction = 'backward')
```

Start: AIC=-2113.69

pointspermatch ~ winspermatch + goalspermatch + goalsagainstpermatch +  
goalmarginpermatch + topfiveleague

Step: AIC=-2113.69

pointspermatch ~ winspermatch + goalspermatch + goalsagainstpermatch +  
topfiveleague

	Df	Sum of Sq	RSS	AIC
<none>			9.6400	-2113.7
- topfiveleague	1	0.0446	9.6846	-2113.2
- goalspermatch	1	0.7904	10.4305	-2073.9
- goalsagainstpermatch	1	3.9802	13.6203	-1932.5
- winspermatch	1	14.4821	24.1221	-1629.6

Call:

```
lm(formula = pointspermatch ~ winspermatch + goalspermatch +  
goalsagainstpermatch + topfiveleague, data = soccer)
```

Coefficients:

(Intercept)	winspermatch	goalspermatch
0.43510	1.38121	0.09647
goalsagainstpermatch	topfiveleagu	top five
-0.09789	0.02570	

```
bestbackward = lm(pointspermatch ~ winspermatch + goalspermatch + goalsagainstpermatch +
  topfiveleague, data = soccer)
```

## Both Selection

```
stepAIC(m_none,
  scope = list(lower = m_none, upper = m_all),
  data = soccer, direction = 'both')
```

Start: AIC=-977.7  
 pointspermatch ~ 1

	Df	Sum of Sq	RSS	AIC
+ winspermatch	1	69.170	14.292	-1911.0
+ goalmarginpermatch	1	54.871	28.590	-1543.5
+ goalspermatch	1	41.898	41.564	-1345.2
+ goalsagainstpermatch	1	33.999	49.463	-1253.0
+ topfiveleague	1	9.140	74.322	-1037.2
<none>			83.462	-977.7

Step: AIC=-1910.99  
 pointspermatch ~ winspermatch

	Df	Sum of Sq	RSS	AIC
+ goalmarginpermatch	1	4.607	9.685	-2115.2
+ goalsagainstpermatch	1	3.740	10.552	-2069.8
+ goalspermatch	1	0.467	13.825	-1926.6
+ topfiveleague	1	0.292	14.000	-1919.9
<none>			14.292	-1911.0
- winspermatch	1	69.170	83.462	-977.7

Step: AIC=-2115.24  
 pointspermatch ~ winspermatch + goalmarginpermatch

	Df	Sum of Sq	RSS	AIC
+ topfiveleague	1	0.0445	9.6402	-2115.7
<none>			9.6847	-2115.2
+ goalspermatch	1	0.0000	9.6846	-2113.2
+ goalsagainstpermatch	1	0.0000	9.6846	-2113.2
- goalmarginpermatch	1	4.6072	14.2919	-1911.0
- winspermatch	1	18.9058	28.5905	-1543.5



Step: AIC=-2115.68

pointspermatch ~ winspermatch + goalmarginpermatch + topfiveleague

	Df	Sum of Sq	RSS	AIC
<none>			9.6402	-2115.7
- topfiveleague	1	0.0445	9.6847	-2115.2
+ goalspermatch	1	0.0002	9.6400	-2113.7
+ goalsagainstpermatch	1	0.0002	9.6400	-2113.7
- goalmarginpermatch	1	4.3594	13.9996	-1919.9
- winspermatch	1	18.4525	28.0927	-1550.8

Call:

```
lm(formula = pointspermatch ~ winspermatch + goalmarginpermatch +  
    topfiveleague, data = soccer)
```

Coefficients:

(Intercept)	winspermatch	goalmarginpermatch
0.43402	1.37907	0.09770
topfiveleague	top five	
0.02556		

```
bestboth = lm(pointspermatch ~ winspermatch + goalmarginpermatch + topfiveleague, data = s
```

## Variable Selection - Mallo's CP

```
soccer_allsub = regsubsets(pointspermatch ~ winspermatch + goalspermatch +  
                             goalsagainstpermatch + goalmarginpermatch + topfiveleague,  
                             data = soccer,  
                             nbest = 1, nvmax = 5)
```

Warning in leaps.setup(x, y, wt = wt, nbest = nbest, nvmax = nvmax, force.in =  
force.in, : 1 linear dependencies found

Reordering variables and trying again:

```
soccer_allsub
```

```

Subset selection object
Call: regsubsets.formula(pointspermatch ~ winspermatch + goalspermatch +
      goalsagainstpermatch + goalmarginpermatch + topfiveleague,
      data = soccer, nbest = 1, nvmax = 5)
5 Variables (and intercept)

              Forced in Forced out
winspermatch      FALSE      FALSE
goalspermatch     FALSE      FALSE
goalsagainstpermatch FALSE      FALSE
topfiveleaguetop five FALSE      FALSE
goalmarginpermatch FALSE      FALSE
1 subsets of each size up to 4
Selection Algorithm: exhaustive

```

```
summary(soccer_allsub)$rsq
```

```
[1] 0.8287614 0.8839632 0.8844959 0.8844979
```

```
summary(soccer_allsub)$which
```

```

(Intercept) winspermatch goalspermatch goalsagainstpermatch
1          TRUE          TRUE          FALSE          FALSE
2          TRUE          TRUE          FALSE          FALSE
3          TRUE          TRUE          FALSE          FALSE
4          TRUE          TRUE          TRUE           TRUE
goalmarginpermatch topfiveleaguetop five
1          FALSE          FALSE
2          TRUE          FALSE
3          TRUE          TRUE
4          FALSE          TRUE

```

```
bestallsubset = lm(pointspermatch ~ winspermatch, data = soccer)
```

## Comparing RMSE after variable selection

RMSE All Subset: 0.1642128

RMSE Best Backward: 0.1348656 - LOWEST

RMSE Best Both, Forward, Lasso: 0.1348667

(they had the same predictors in it)

CONCLUSION:

- best backward has lowest rmse so better
- predictors: wins per match, top five league, goals per match, goals against per match

```
rmse(bestallsubset)
```

```
[1] 0.1642128
```

```
rmse(bestbackward)
```

```
[1] 0.1348656
```

```
rmse(bestboth)
```

```
[1] 0.1348667
```

```
rmse(bestforward)
```

```
[1] 0.1348667
```

```
rmse(bestlasso)
```

```
[1] 0.1348667
```

##possibly delete – Comparing Models: RMSE

RMSE Model 1: 0.1348656 RMSE Model 2: 0.1348656 RMSE Model 3: 0.1348667 RMSE  
Model 4: 0.1349185

```
rmse(model1)
```

```
[1] 0.1348656
```

```
rmse(model2)
```

```
[1] 0.1348656
```

```
rmse(model3)
```

```
[1] 0.1348667
```

```
rmse(model4)
```

```
[1] 0.1349185
```

## Results

## Conclusion